

Comparison of Underlying Data Structures for Distributed Ledgers

Sebastian Sanchez and Nicolas Cardozo

Abstract— Cryptocurrencies, especially Bitcoin, have achieved notable financial and public recognition in the recent past but they have had problems due to poor scaling. The issues are in clear sight of developers, with different answers to the challenges being proposed, avid discussion within the community, and coding of potential solutions. The issues have also prompted the proposal of different underlying data structures for distributed ledgers different from the blockchain. The aim of this thesis is to address the optimality of these data structures, by determining if and how they guarantee the fundamental properties of blockchains and if they can respond to some of the technical challenges identified for blockchains.

Index Terms— blockchain, block-lattice, data structure, distributed ledger, tangle, trade-off

I. CONTEXT

BLOCKCHAIN is the technology behind Bitcoin. This technology provides an append-only data store of transactions replicated between peers. A blockchain implements a distributed ledger, which can verify and store any kind of transactions. Many institutions around the world are exploring the applications of distributed ledgers in areas such as supply chain, electronic health records, voting, energy supply, ownership management, and protecting civil infrastructure [1].

The distributed ledger industry is in its early stages of development. There are different kinds of limitations already identified for blockchain as a solution for these systems. These limitations include those related to technical issues with the underlying algorithms and data structures used, ongoing industry thefts and scandals, public perception, government regulation, and mainstream adoption of the technology.

The Bitcoin blockchain prompted a rise in the research of distributed ledger technology. Researchers around the world have proposed different data structures that can underlie these ledgers and which, they claim, can respond to some of blockchains potential limitations. The objective of this thesis is to determine whether there exist viable alternatives to the blockchain, from the perspective of the underlying data structures used to build distributed ledgers.

II. RESEARCH PROBLEM

With the recent increase in the popularity of Bitcoin, blockchain has become a prime keyword in the computer science market. The blockchain has risen to become the most popular implementation of distributed ledgers and distributed applications in general. Nonetheless, one of the reasons to use blockchain implementations in distributed systems is centered around the properties it offers in terms of security (in the broader and popular use of the term). Blockchain is designed around the concept of trustless systems, which can be of interest for some application domains. However, implementations of distributed ledgers that use different data structures have started to appear. This thesis will try to answer whether blockchain's status is justified or if other implementations can replace the blockchain as the most popular implementation of distributed ledgers.

A. Fundamental Properties

Satoshi Nakamoto created the Bitcoin blockchain as peer-to-peer network that could serve as a solution to double-spending; a problem that hindered the development of decentralized digital currency for decades [2]. In the Bitcoin whitepaper, Nakamoto explained,

“A purely peer-to-peer version of electronic cash would allow online payments to be sent directly from one party to another without going through a financial institution. Digital signatures provide part of the solution, but the main benefits are lost if a trusted third party is still required to prevent double-spending. We propose a solution to the double-spending problem using a peer-to-peer network. The network timestamps transactions by hashing them into an ongoing chain of hash-based proof-of-work, forming a record that cannot be changed without redoing the proof-of-work. The longest chain not only serves as proof of the sequence of events witnessed, but proof that it came from the largest pool of CPU power. As long as a majority of CPU power is controlled by nodes that are not cooperating to attack the network, they'll generate the longest chain and outpace attackers. The network itself requires minimal structure. Messages are broadcast on a best effort basis, and nodes can leave

and rejoin the network at will, accepting the longest proof-of-work chain as proof of what happened while they were gone.” [2, p. 1]

After carefully studying blockchain implementations, Xu et al. [3] identified the fundamental properties supported by this family of technological solutions. Any solution targeting the blockchain technology should guarantee the five properties below:

1) *Immutability*

In the blockchain context, immutability refers to the fact no one is able to change the data that is contained in a committed transaction. If a malicious user attempts to alter the information in a transaction, this change will cause a change in the hash of the block, causing the protocol to register it as an inconsistency in the chain.

2) *Non-repudiation*

Blockchain users sign the transactions they initiate with digital signatures, which means that they cannot dispute the authorship or validity of such a transaction.

3) *Integrity*

The blockchain guarantees data integrity, which means it maintains and assures the accuracy and consistency of data. There are mechanisms in place by which all parties in the system can reach a consensus on the accepted truth. The Bitcoin blockchain achieves consensus via an economic measure called Proof of Work (PoW).

4) *Transparency*

On a blockchain, the identity of a user is concealed using cryptography so that linking public addresses to individual users is difficult to achieve. The transparency of a blockchain comes from the fact that the holdings and transactions of each public address are open to viewing. This level of transparency has not existed within financial systems and adds a degree of accountability not existing to date.

5) *Equal Rights*

All of the participants of a blockchain network have equal rights, and there are no privileged users because they all have the same ability to access and manipulate the blockchain.

For a solution to be comparable to a blockchain it must guarantee these five fundamental properties or give a reasoned explanation as to why it does not do so. Researchers designing different solutions similar to blockchains must keep these five properties in mind.

B. *Technical Challenges*

A number of technical challenges related to the blockchain have been identified. Developers have a clear view of these challenges, which has led to discussions, coding of possible solutions and proposals of different answers to these challenges. Some researchers believe the Bitcoin blockchain will remain as the solution with the biggest market value, while others are building new solutions that they claim could overcome these

challenges and become the next steps in the evolution of this technology [4], [5].

Melanie Swan, Founder of the Institute for Blockchain Studies and author of the foremost reference regarding blockchain: *Blockchain: Blueprint for a New Economy* [6] identifies a number of technical challenges related to the blockchain. These are explained below.

1) *Throughput*

Throughput is the amount of transactions that can be processed per second and it is an issue for the Bitcoin network. The on-chain transaction processing capacity of the Bitcoin network is limited by the average block creation time of 10 minutes and the block size limit of 1 MB. These jointly constrain the network’s throughput. This network is processing only one transaction per second (tps) and has a theoretical maximum of 7 tps.

Other transaction processing networks have much higher throughputs such as Visa, which supports 2,000 tps, Twitter, which supports 5,000 tps, and advertising networks, which support more than 100,000 tps. Blockchain’s low throughput is a major concern for enterprises that depend on high-performance legacy transaction processing systems. To be ready for pervasive use, distributed ledgers would need to be able to process more transactions per second.

2) *Latency*

Latency is the time it takes from the creation of a transaction until the initial confirmation of it being accepted by the network. Latency is an issue for the Bitcoin network. Each Bitcoin transaction block takes 10 minutes to process, meaning it takes at least 10 minutes for a transaction to be confirmed as accepted. As a comparison metric, VISA takes seconds at most to process a transaction. Latency is an issue that needs to be solved if this technology is to compete with payment methods currently in use.

C. *Justification*

Taking into account the fundamental properties of blockchain technology and two of its technical challenges, we state that the implementation of distributed ledgers by making use of a blockchain is insufficient. In order to truly realize a distributed ledger, new data structures and algorithms need to be designed. Hereinafter, we present an evaluation of different alternatives for the implementation of distributed ledgers from the perspective of the underlying data structure used to manage them.

Taking into account the technical challenges, different data structures have been proposed to replace the blockchain as the data structure used to implement distributed ledgers. The problem we address with this thesis is that of the advantages and disadvantages of the underlying data structures used to implement distributed ledgers. We study different data structures to determine how they compare in terms of the degree to which (1) they guarantee the fundamental properties of blockchains, and (2) they respond to some of the technical challenges that have been identified for this technology.

In particular, to address this problem we will follow these four steps:

1. Implement distributed data structures in a way that allows for comparison in terms of the two technical challenges identified for this technology.
2. Qualitatively evaluate and compare the data structures in terms of the five fundamental properties of blockchains.
3. Quantitatively evaluate and compare the data structures in terms of the two technical challenges.
4. Classify the different data structures according to their impact on the two technical challenges identified for this technology.

III. STATE OF THE ART

Before we present our evaluation study, we describe the current state of the art in blockchain technology. This section is focused in the implementation of blockchains, rather than their use in other systems.

A. Methodology

The methodology used to gather the state of the art for this research follows the systematic review process described by Petersen et al. [7], and is informed by guidelines for a systematic literature review described by Kitchenham [8]. The steps followed are (1) the definition of the research area, (2) the search for papers, and (3) the screening for relevant papers.

1) Definition of the Research Area

The first stage of the process is the definition of the research area. The research area is defined by two topics:

1. *Evaluation of the blockchain* - rigorous scientific evaluation and validation of the blockchain as a distributed ledger.
2. *Distributed Ledgers and their Underlying Data Structures* – alternative data structures that can underlie distributed ledgers.

2) Search for Papers

The second stage of the systematic review is the search for papers. The search for papers includes scientific databases and technical text books, but it also needs to include working papers (e.g., white papers), considering the fact that the blockchain industry is still in the early stages of development. The chosen databases for the search are (1) ACM Digital Library, (2) IEEE Xplore, and (3) Springer Link, because of the high quality of the papers published in these scientific databases. The terms used in the search string are different according to the two topics that were described before, and they are shown below:

1. *Evaluation of the blockchain*: *evaluati** and *blockchain**
2. *Distributed Ledgers and their Underlying Data Structures*
 - a. *Distributed Ledger Technology*: *distributed* and *ledger** and *technolog**

- b. *Underlying Data Structures for Distributed Ledgers*: (*blockchain** and *alternative**) or (*blockchain** and *substitute**)

The interest in blockchain technology has been increasing since the idea was first introduced in January 2009. In 2015, Yli-Huuma et al. [9] did a systematic mapping study to describe the state of research on blockchain technology and found that the cumulative number of papers written about blockchain from a computer science perspective increased from 2012 until 2015, when their study was published. This trend can be seen in Figure 1. This group of authors provides some recommendations on future research directions of blockchain technology. One of their recommendations was for more studies to be conducted on the scalability issues of blockchain, which this thesis does by analyzing throughput and latency.




Figure 1. Increasing number of publications related to blockchain per year. Taken from [9].

Dabbagh, Sookhak and Safa [10] conducted a bibliometric study of blockchain-related publications indexed by Web of Science from 2013 to 2018. Figure 2 shows a generally growing trend in the number of blockchain papers indexed by Web of Science per year.




Figure 2. Number of blockchain papers indexed by Web of Science per year. Taken from [10].

3) Screening of Papers for Inclusion and Exclusion

The last stage of the process is the screening of relevant papers. In this stage, some of the papers that were found had to be discarded because they were not necessarily related to the research topics. This process is described in detail in the following section.

4) Search and Selection Results

The search and selection results for each of the three research topics are presented below.

1. *Evaluation of the blockchain*: 6 papers were found after the search. 2 were selected and 4 were excluded. Some of the papers that were excluded evaluated blockchains in specific settings, such as business process execution [11] and cloud [12], but what was needed was a general evaluation of blockchain.
2. *Distributed Ledgers and their Underlying Data Structures*
 - a. *Distributed Ledger Technology*: 7 papers were found after the search. 5 were selected and 2 were excluded. The papers that were excluded were two literary reviews of distributed ledgers from a business perspective, and therefore were not related to the research topic.
 - b. *Underlying Data Structures for Distributed Ledgers*: 7 papers were found after the search. 4 were selected and 3 were excluded. The papers that were excluded presented solutions that made use of blockchains, which meant they did not present other alternatives to this data structure.

B. Evaluation of the blockchain

It is important to have an understanding of what blockchains can and cannot do. The following describes evaluations of blockchains aimed at answering these questions.

Xu et al. [3] propose a taxonomy that enables the classification and comparison of blockchains. Their taxonomy is informed by academic literature, books, government and technical reports, documents of industrial blockchain products, and developer forums and wikis. Additionally, they use an investigation for the Australian government of the use of blockchains in various use cases and their experience from implementing proof-of-concept blockchain-based systems. This taxonomy aims to assist with the design and assessment of the impact of blockchains on software architectures and helps with considerations about the quality attributes of blockchain-based systems. The authors analyze architectural design issues for blockchain-based systems in terms of three categories: the level of (de)centralization, the support for client storage and computation, and the blockchain infrastructural configuration. It is worth remembering that this group of researchers were responsible for identifying the five fundamental properties of blockchains.

Ahn Dihn et al. [13] argue that it is important to understand what a blockchain can offer, especially with respect to its data processing capabilities. This group of researchers created BLOCKBENCH, a benchmarking framework that helps understand the performance of private blockchains against data processing workloads. They conduct a comprehensive evaluation of three blockchain systems based on BLOCKBENCH, namely Ethereum, Parity, and Hyperledger Fabric. Their results demonstrated several trade-offs in the

design space, as well as big performance gaps between blockchain and database systems.

According to “Stack Overflow’s annual Developer Survey for 2019” [14], which “is the largest and most comprehensive survey of people who code around the world” [14], and which was taken by nearly 90,000 developers from over 170 countries; when asked what they primarily believe about blockchain technology, respondents of the survey are largely optimistic about its broad usefulness. However, this optimism is largely concentrated among young, less experienced developers. The more experienced a respondent is, the more likely he is to say blockchain technology is an irresponsible use of resources.

The papers described in this section present a general view of the capabilities and limitations of blockchains. They were essential in defining the properties and metrics that are analyzed and measured in the experiments conducted in this research.

IV. DISTRIBUTED LEDGERS AND THEIR UNDERLYING DATA STRUCTURES

In this section we present three data structures used to implement distributed ledgers. We focus on the definition of the structures in this section. The following section presents our implementation for each of them.

Before going further, let us properly define ledgers and distributed ledgers. A *ledger* is an account book of final entry, in which business transactions are recorded. A *distributed ledger* is a consensus of replicated, shared, and synchronized digital data where there is no central administrator or centralized data storage [1]. To ensure replication across nodes is undertaken, distributed ledgers require a peer-to-peer network and a consensus algorithm [15]. Distributed ledgers make use of data structures to store the transactions.

A. The blockchain

When Satoshi Nakamoto set the Bitcoin blockchain into motion in 2009, he introduced the concept of a proof of work-based blockchain to allow an agreement on the order of transactions [2]. The blockchain is the first credible solution to the double-spending problem, which for decades hindered the development of decentralized digital currency.

As a data structure, the blockchain, pictured in Figure 3, is an ordered list of blocks where each block contains a list of transactions. Each block in the blockchain is “chained” back to the previous block by containing a hash of the previous block. This way, the historical transactions in the blockchain may not be deleted or altered without invalidating the chain of hashes.




Figure 3. The blockchain. Image taken from [2].

Nakamoto explains that the steps to run the network are the following:

- 1) New transactions are broadcast to all nodes.
- 2) Each node collects new transactions into a block.
- 3) Each node works on finding a difficult proof-of-work for its block.
- 4) When a node finds a proof-of-work, it broadcasts the block to all nodes.
- 5) Nodes accept the block only if all transactions in it are valid and not already spent.
- 6) Nodes express their acceptance of the block by working on creating the next block in the chain, using the hash of the accepted block as the previous hash.

Nodes always consider the longest chain to be the correct one and will keep working on extending it.

B. The tangle

IOTA is a new a cryptocurrency designed for the Internet-of-Things industry. Sergei Popov introduces the main feature of this cryptocurrency: the *tangle* [4]. The tangle, pictured in Figure 4, is a directed acyclic graph for storing transactions and, according to the group behind IOTA, “succeeds the blockchain as its next evolutionary step” [4, p. 1].

To clarify the terminology, vertices are the transactions represented on the tangle graph and nodes are entities that issue and validate transactions. The transactions issued by nodes are the vertex set of the tangle graph and the edge set is obtained in the following way: when a new transaction arrives, it must approve two previous transactions. These approvals are represented by directed edges from the new transactions to the other two transactions. If there is not a directed edge between transaction A and transaction B, but there is a directed path of length at least two from A to B, it is said that A indirectly approves B.

The main idea behind the tangle is that to issue a transaction, users must work to approve other transactions. In this way, users who issue a transaction are contributing to the network’s security. Nodes check if transactions are conflicting. If a node finds that a transaction is in conflict with the tangle history, the node will not approve the conflicting transaction. As a transaction receives additional approvals, it is accepted by the system with a higher level of confidence.




Figure 4. The tangle. Image taken from [16].

Popov explains that to issue a transaction, a node does the following:

- 1) The node chooses two other transactions to approve according to an algorithm. In general, these two transactions may coincide.
- 2) The node checks if the two transactions are conflicting and does not approve conflicting transactions.

- 3) For a node to issue a valid transaction, the node must solve a cryptographic puzzle similar to those in the Bitcoin blockchain. This is achieved by finding a nonce (an arbitrary number used only once in a cryptographic communication) such that the hash of that nonce concatenated with some data from the approved transaction has a particular form.

The tangle may contain conflicting transactions. However, in this case, the nodes need to decide which transactions will become orphaned. Orphaned transactions are not approved by incoming transactions and therefore do not achieve a high level of confirmation confidence which is a measure of a transaction’s level of acceptance by the rest of the tangle.

In the same way as a block in a blockchain cannot be modified without invalidating all the subsequent blocks in the chain, a transaction in the tangle cannot be modified without invalidating all the transactions that approve that transaction.

C. The block-lattice

The block-lattice, picture in Figure 5, is the data structure used to store the information for the Nano cryptocurrency [5]. In the block-lattice, each account has its own blockchain (account-chain) which stores the account’s transaction history. Every node in the network stores a ledger composed of its account-chain and a copy of the account-chains of all the other nodes in the network.




Figure 5. The block-lattice. Image taken from [5].

Each account-chain can only be updated by the account’s owner, which allows each account-chain to be updated asynchronously to the rest of the block-lattice. Every transfer of funds requires the creation of a send transaction (S) on the sender’s account-chain and the creation of a receive transaction (R) on a receiver’s account-chain. The transfer is completed only when both transactions are acknowledged and accepted by the network after being broadcast by the respective account-chain owners.

V. LEDGER DATA STRUCTURE IMPLEMENTATIONS

This section shows the implementation of the three data structures underlying distributed ledgers introduced in the previous section, namely the blockchain, the tangle, and the

block-lattice. To implement the data structures, we use the Go programming language¹ in order to take advantage of its distributed programming characteristics.

A. Go in a Nutshell

After studying the advantages and disadvantages of different programming languages, the language selected for the implementation of the data structures is Go [17]. Go, also referred to as Golang, is a statically typed, compiled programming language designed at Google. Some of its features include memory safety, garbage collection, structural typing, and CSP-style concurrency. This language is selected because its concurrency mechanisms make it easy to write programs that get the most out of multicore and networked machines. It has two main built-in facilities for writing concurrent and distributed programs, goroutines and channels, which are described in more detail below.

1) Goroutines

Goroutines [18] are functions that run concurrently with other functions. They can be thought of as light weight threads as the cost of creating a goroutine is much less than the cost of creating a thread. Goroutines are a few KB in stack size and the stack can grow and shrink according to the needs of the application whereas, with threads, the stack size has to be specified and is fixed.

The language abstraction `go f(x, y, z)` starts a new goroutine running the `f(x, y, z)` function. The evaluation of `f`, `x`, `y`, and `z` happens in the currently executing goroutine (the one calling the `go` abstraction), and the execution of `f` happens in the new goroutine.

The code snippet shown in Figure 6 presents a go package with two functions `main` and `say`. Note in this snippet that the first call to the `say` function takes place in an independent go routine, while the second call happens in the currently executing go routine. The output when executing the main function is presented in Figure 7.

```
package main

import (
    "fmt"
    "time"
)

func say(s string) {
    for i := 0; i < 5; i++ {
        time.Sleep(100 * time.Millisecond)
        fmt.Println(s)
    }
}

func main() {
    go say("hello")
    say("world")
}
```

Figure 6. Code snippet for Goroutines

```
hello
world
world
hello
hello
world
world
hello
hello
world

Program exited.
```

Figure 7. Output for Goroutines code snippet in Figure 6.

Notice the interleaving between the prints of “hello” and “world”. This interleaving reflects the goroutines running concurrently.

2) Channels

Goroutines communicate using channels [19]. Channels prevent race conditions from happening when accessing shared memory using goroutines, and they allow goroutines to synchronize without explicit locks or condition variables. Channels can be conceived as communication buffers allowing to read and write multiple values.

Channels send and receive values with the channel operator: `<-` whose functioning is shown in Figure 8.

```
//create a channel
ch := make(chan int)

// Send v to channel ch
ch <- v

// Receive from ch, and assign value to v
v := <-ch
```

Figure 8. How Go Channels work.

The code snippet shown in Figure 9 sums the numbers in a slice (a dynamically-sized array in Go). It distributes the work between two goroutines to provide data parallelism. Once both goroutines have completed their computation, it calculates the final result. The output when executing the main function is presented in Figure 10.

¹ <https://golang.org/>

```

package main

import "fmt"

func sum(s []int, c chan int) {
    sum := 0
    for _, v := range s {
        sum += v
    }
    c <- sum // send sum to c
}

func main() {
    s := []int{7, 2, 8, -9, 4, 0}

    c := make(chan int)
    go sum(s[:len(s)/2], c)
    go sum(s[len(s)/2:], c)
    x, y := <-c, <-c // receive from c

    fmt.Println(x, y, x+y)
}

```

Figure 9. Code snippet for Go Channels.

```

-5 17 12
Program exited.

```

Figure 10. Output for Go Channels code snippet in Figure 9.

In the setting of this thesis, read and write operations have to be performed in an asynchronous manner because nodes in the network have to be able to receive and send the shared data structures. Goroutines and channels help to achieve this effectively.

B. Peer-to-Peer Communication

In a distributed setting, nodes in the network must communicate with each other without the need for a central server. To do this, we use a peer-to-peer (p2p) networking stack called noise [20]. Noise is an easy-to-use networking stack for developing decentralized applications written in Go and developed by the Perlin Team [21].

When a node runs a Go program using noise, it acts as a peer to which other nodes can connect, and nodes can connect and disconnect from the network without causing a disruption to the communication. Noise also allows nodes to broadcast messages over their network of peers.

C. Data Structure Design

The blocks in the IOTA tangle and the Nano block-lattice have a granularity of one transaction. In order to enable a comparison, the data structures implemented in this thesis are comprised of items which record one transaction each. This marks a clear difference from the way in which the Bitcoin blockchain is implemented since the blocks in that system hold the data for various transactions.

1) Blockchain Implementation

Our Go implementation of the blockchain uses the implementation presented by Coral Health [22] as a starting point. The blockchain consists of a series of blocks which each

have an index, a timestamp, a hash, and the hash of the previous block as shown in Figure 11.

```

// Block represents each transaction in the
// blockchain
type Block struct {
    Index      int
    Timestamp  string
    Transaction string
    Hash       string
    PrevHash   string
    Signature  string
    TimeSent   time.Time
}

// Blockchain is a series of validated Blocks
type Blockchain struct {
    Blocks    []Block
    State     map[string]int
    Difficulty int
}

```

Figure 11. Basic code for the blockchain.

There are a number of specific functions for the blockchain.

`isBlockValid` checks that the blockchain is consistent. It does this by: (1) checking that the index of a new block is one more than the index of the previous block, (2) checking that the reference to the previous block is correct, and (3) verifying the hash for a new block. Figure 12 shows the code for this function.

```

// Make sure block is valid by checking index,
// comparing the hash of the previous block,
// and verifying hash
func isBlockValid(newBlock, oldBlock
network.Block) bool {
    if oldBlock.Index+1 != newBlock.Index {
        return false
    }

    if oldBlock.Hash != newBlock.PrevHash {
        return false
    }

    if calculateHash(newBlock) != newBlock.Hash {
        return false
    }

    return true
}

```

Figure 12. `isBlockValid` function.

`calculateHash` uses sha256 to hash a block's raw data and return a string which is the hash for the block. Figure 13 shows the code for this function.

```

// SHA256 hashing
func calculateHash(block network.Block) string {
    record := strconv.Itoa(block.Index) +
        block.Timestamp + block.Transaction +
        block.PrevHash + block.Signature
    h := sha256.New()
    h.Write([]byte(record))
    hashed := h.Sum(nil)
    return hex.EncodeToString(hashed)
}

```

Figure 13. `calculateHash` function.

generateBlock creates a new block to be added to the blockchain with the necessary transaction information inside it. Figure 14 shows the code for this function.

```
// create a new block using previous block's
// hash
func generateBlock(oldBlock network.Block,
Transaction string, address string, timeSent
time.Time) network.Block {

    var newBlock network.Block

    t := time.Now()
    newBlock.Timestamp = time.Unix(0,
t.UnixNano()).String()

    newBlock.Index = oldBlock.Index + 1
    newBlock.Transaction = Transaction
    newBlock.PrevHash = oldBlock.Hash
    newBlock.Signature = address
    newBlock.TimeSent = timeSent

    newBlock.Hash = calculateHash(newBlock)

    return newBlock
}
```

Figure 14. generateBlock function.

2) Tangle implementation

The company in charge of introducing the tangle created a visual simulation written in React and D3.js [16]. This visualization is used as a guide for the implementation of the tangle in Go.

The tangle is defined as a Directed Acyclic Graph (DAG), where nodes represent blocks, each containing exactly one transaction, and edges represent verification links between transactions. Each transaction is identified by an index, as shown in Figure 15.

```
type Transaction struct {
    Index      int
    Operation  string
    TimeInt    int64
    TimeString string
    Weight     int
    CumWeight  int
    Signature  string
    Hash       string
    HashApp1   string
    HashApp2   string
    TimeSent   time.Time
}

type Link struct {
    Target int
    Source int
}

// Tangle is a DAG of Transactions
type DAG struct {
    Transactions []Transaction
    Links        []Link
    Lambda      float64
    Alpha        float32
    H            int64
    TipSelection string
    State        map[string]int
}
```

Figure 15. Basic code for the tangle.

There are a number of specific functions for the tangle.

generateTransaction creates a new transaction to be added to the tangle with the necessary transaction information inside of it. Figure 16 shows the code for this function.

```
// create a new Transaction using previous
// Transactions index
func generateTransaction(lastTransaction
network.Transaction, Operation string, address
string, timeSent time.Time) network.Transaction
{

    var newTransaction network.Transaction

    newTransaction.Index =
lastTransaction.Index+1
    newTransaction.Operation = Operation

    now := time.Now()
    newTransaction.TimeInt = now.UnixNano() /
1000000
    newTransaction.TimeString = time.Unix(0,
now.UnixNano()).String()
    newTransaction.Weight = 1
    newTransaction.Signature = address
    newTransaction.Hash =
calculateHash(newTransaction)

    newTransaction.TimeSent = timeSent

    return newTransaction
}
```

Figure 16. generateTransaction function.

generateLink creates a new link, from a source (approving) transaction to a target (approved) transaction, to be added to the tangle. Figure 17 shows the code for this function.

```
func generateLink(target network.Transaction,
source network.Transaction) network.Link {

    var newLink network.Link

    newLink.Target = target.Index
    newLink.Source = source.Index

    return newLink
}
```

Figure 17. generateLink function.

3) Block-lattice Implementation

The whitepaper “Nano: A Feeless Distributed Cryptocurrency Network” [5] is used as a starting point for the implementation of the block-lattice with Go.

The block-lattice is made of account-chains that store transactions. These transactions have an index, a transaction type i.e., send and receive, the *amount* being sent/received in the transaction, and the resulting *balance* in the account after the transaction has been completed. The main structure of a transaction block is shown in Figure 18.


```
// Cube represents each transaction
// in the block-lattice
type Cube struct {
    Index      int
    Balance    int
    Type       string
    Amount     int
    Hash       string
    Source     string
    Previous   string
    Signature  string
    TimeSent   time.Time
    Sender     string
    Receiver   string
}
```

Figure 18. Basic code for the block-lattice

A specific function for the block-lattice is `generateCube`. This function creates a new cube to be added to the block-lattice with the necessary information inside of it. Notice that if it's a send transaction the amount being transferred is deducted from the account's balance, and if it is a receive transaction the amount being transferred is added to the account's balance. Figure 19 shows the code for this function.

```
// create a new cube
func generateCube(oldCube network.Cube,
    typeOfTransaction string, amount int, signature
    string, hash string, timeSent time.Time, sender
    string, receiver string) network.Cube {

    var newCube network.Cube

    newCube.Index = oldCube.Index + 1
    newCube.Previous = oldCube.Hash

    if typeOfTransaction == "send" {
        newCube.Balance = oldCube.Balance - amount
    } else {
        newCube.Balance = oldCube.Balance + amount
    }

    newCube.Type = typeOfTransaction
    newCube.Amount = amount

    if hash == " " {
        newCube.Hash = calculateHash(newCube)
    } else {
        newCube.Hash = hash
    }

    newCube.Signature = signature
    newCube.TimeSent = timeSent

    newCube.Sender = sender
    newCube.Receiver = receiver

    return newCube
}
```

Figure 19. `generateCube` function.

VI. EVALUATION OF DATA STRUCTURES

The next steps after implementing the data structures is to evaluate and classify them with respect to the characteristics presented in Section II. The purpose of this classification is to help developers and practitioners decide which data structure to use depending on the specific characteristics of their application

domain and according to the fundamental properties they wish to guarantee and the quality attributes they want to prioritize.

A. Design of Evaluation

1) Quantitative Evaluation and Comparison

The key metrics that are evaluated for the blockchain, the tangle, and the block-lattice are throughput and latency. To measure these metrics, transactions are submitted to the network by a node which stores the transactions in its copy of the data structure and then shares this information with its peers.

Each node is deployed in a different virtual machine instance. The instances are setup in a Google Cloud Platform cluster, where each instance is an `n1-standard-1` instance that has one virtual CPU, 3.75GB of RAM, 10 GB of hard drive, and runs a Linux operating System. The experiments that are conducted in this research resemble the experiments conducted by Ahn Dihn et al. [13] in their comparative study of three blockchain systems (namely Ethereum, Parity, and Hyperledger Fabric).

To measure throughput, one node submits as many transactions as it can for one second. Each throughput experiment is repeated five times and the median and standard deviation for the five runs are reported. To measure latency one node submits transactions to the network and, for each transaction, the receiving node(s) subtract the initiation timestamp from the completion timestamp. The median and standard deviation are reported. The completion timestamp for transactions using the block-lattice is taken after the receive transaction is completed.

2) Analysis of Fundamental Properties

To determine if and how the three data structures satisfy the fundamental properties of distributed ledgers, the documentation for each of these data structures is analyzed. The result of this analysis is a description of the manner in which each data structure guarantees each property. In case a data structure does not guarantee a certain property there is an explanation of why this is the case.

3) Classification of Data Structures

After gathering and analyzing the quantitative information for each of the data structures, we classify them according to their impact on the quality attributes of throughput and latency. The method of classification that is used mirrors the method used by Xu et al. in their "Taxonomy of Blockchain-Based Systems" [3]. Xu et al. use a classification scheme to describe the impact of architectural decisions over a group of quality attributes by stating whether each option was less favorable (\oplus), neutral ($\oplus\oplus$), or more favorable ($\oplus\oplus\oplus$) with respect to each attribute in comparison established between the different options that are available.

VII. RESULTS OF EVALUATION

A. Results of Quantitative Analysis

After conducting the quantitative analysis, the following conclusions can be made regarding the degree to which the

different data structures satisfy important quality attributes of distributed ledgers.

1) Throughput

The data structure with the best throughput by a large margin is the block-lattice. Figure 20 shows that the difference in throughput is very large between the block-lattice and the other two data structures. It also shows that the blockchain has a better throughput than the tangle which has the worst throughput of the three.

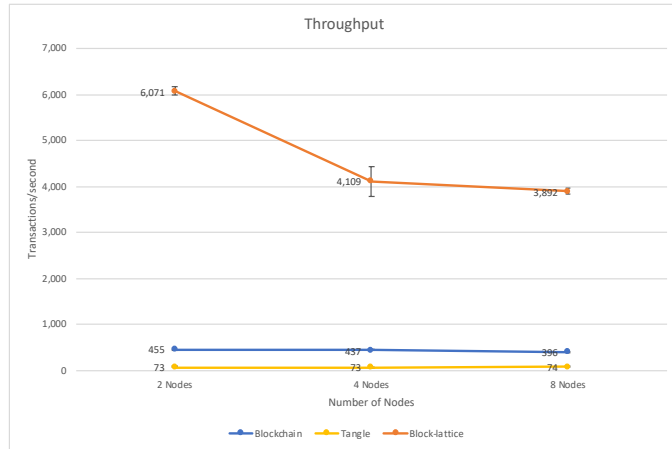


Figure 20. Throughput for the three data structures.

2) Latency

The latency of a transaction is the time spent from the moment of its submission until it is completed. This time may vary depending on the size of the data structure, reason why it is measured for different transactions as each data structure continues to grow (transaction number: 50, 100, 200, 400, and 800).

The data structure with the best latency is the blockchain. Figure 21 (network with 8 nodes), Figure 22 (network with 4 nodes), and Figure 23 (network with 2 nodes) show that from transaction number 50 onwards, the blockchain has a better latency than the block-lattice and the tangle. The latency for the tangle for transaction number 100 and above is in the order of seconds and massively exceeds that of the other two data structures. The experimental data shows that median latency is more-or-less independent of the network sizes that were chosen.




Figure 21. Latency for the three data structures on a network with 8 nodes.




Figure 22. Latency for the three data structures on a network with 4 nodes.




Figure 23. Latency for the three data structures on a network with 2 nodes.

B. Results of Qualitative Analysis

1) The blockchain

In Section II - Research Problem, the five fundamental properties of distributed ledgers are defined together an explanation of the manner in which blockchains guarantee these properties. However, a more detailed study of blockchains opens room to make further comments about the degree to which the blockchain guarantees equal rights for participants in the network.

a) Equal Rights

In the Bitcoin blockchain system, there are two distinct types of participants, those who issue transactions, and those who approve transactions, known as miners. According to Popov, the mathematician behind the tangle, “the design of this system creates unavoidable discrimination of some participants ...”[4, p. 1]. LeMahieu, the researcher who proposes the block-lattice, explains this in more detail:

“Bitcoin, and other cryptocurrencies, function by achieving consensus on their global ledgers in order to verify legitimate transactions while resisting malicious actors. Bitcoin achieves consensus via an economic measure called Proof of Work (PoW). In a PoW system participants compete to

compute a number, called a nonce, such that the hash of the entire block is in a target range. This valid range is inversely proportional to the cumulative computation power of the entire Bitcoin network in order to maintain a consistent average time taken to find a valid nonce. The finder of a valid nonce is then allowed to add the block to the blockchain; therefore, those who exhaust more computational resources to compute a nonce play a greater role in the state of the blockchain. PoW provides resistance against a Sybil attack, where an entity behaves as multiple entities to gain additional power in a decentralized system, and also greatly reduces race conditions that inherently exist while accessing a global data-structure.”[5, p. 1]

Since the other fundamental properties are defined for and from the blockchain, this data structure satisfies them all.

2) The tangle

a) Immutability

The tangle guarantees that no one is able to change the data that is contained in a committed transaction by using hashes. If a malicious user attempts to alter the information in a transaction, this will cause a change in the hashes of the tangle, causing the protocol to register it as an inconsistency in the data structure. “For a node to issue a valid transaction, the node must solve a cryptographic puzzle similar to those in the Bitcoin blockchain. This is achieved by finding a nonce such that the hash of that nonce concatenated with some data from the approved transaction has a particular form” [4, p. 3].

b) Non-repudiation

Tangle users sign the transactions they initiate with digital signatures, which means that they cannot dispute the authorship or validity of such a transaction.

c) Integrity

The tangle guarantees data integrity, which means it maintains and assures the accuracy and consistency of data. There are mechanisms in place by which all parties in the system can reach a consensus on the accepted truth. When using the tangle, consensus is achieved via a cumulative Proof of Work (PoW) of stacked transactions.

d) Transparency

On a tangle, the identity of a user is concealed using cryptography so that linking public addresses to individual users is difficult to achieve. The transparency of a tangle comes from the fact that the transactions of each public address are open to viewing.

e) Equal Rights

All of the participants of a tangle network have equal rights, and there are no privileged users because they all have the same ability to access and manipulate the tangle. With respect to the approval of transactions, “the main idea of the tangle is the following: to issue a transaction, users must work to approve

other transactions. Therefore, users who issue a transaction are contributing to the network's security.”[4, p. 2]

3) The block-lattice

a) Immutability

According to the Nano whitepaper, the cryptocurrency which uses the block-lattice architecture, “a send block is immutable once confirmed. Once broadcasted to the network, funds are immediately deducted from the balance of the sender’s account and wait as pending until the receiving party signs a block to accept these funds. Pending funds should not be considered awaiting confirmation, as they are as good as spent from the sender’s account and the sender cannot revoke the transaction” [5, p. 3].

b) Non-repudiation

Block-lattice users sign the transactions they initiate with digital signatures, which means that they cannot dispute the authorship or validity of such a transaction.

c) Integrity

The block-lattice guarantees data integrity, which means it maintains and assures the accuracy and consistency of data. There are mechanisms in place by which all parties in the system can reach a consensus on the accepted truth. When using the block-lattice, consensus is achieved via a balanced-weighted vote on conflicting transactions.

d) Transparency

On a block-lattice, the identity of a user is concealed using cryptography so that linking public addresses to individual users is difficult to achieve. The transparency of a block-lattice comes from the fact that the holdings and transactions of each public address are open to viewing.

e) Equal Rights

All of the participants of a block-lattice network have equal rights, and there are no privileged users because they all have the same ability to access and manipulate their own account-chain. Regarding transaction verification, “each individual user provides the computational power for the verification of their own transactions, meaning entire network is not required to update the overall ledger together in massive blocks.” [23]

C. Classification of Data Structures

After gathering and analyzing the quantitative information about the two metrics for each of the data structures, they are classified as Table I.

TABLE I. CLASSIFICATION ACCORDING TO IMPACT ON QUALITY ATTRIBUTES.

Design Decision - Data Structure		
(: Less favorable, :: Neutral, ::: More favorable)		
Option	Impact	
	Quality Attributes	
	Throughput	Latency
Blockchain	⊕	⊕⊕⊕
Tangle	⊕	⊕
Block-lattice	⊕⊕⊕	⊕⊕

VIII. IMPORTANCE OF RESULTS

A. Importance of Results of Quantitative Analysis

The quantitative analysis demonstrates there is a trade-off between throughput and latency.

If users wish to prioritize throughput, the best option is the block-lattice. The reason for the high throughput for the block-lattice is believed to be ease with which users can update their own account-chain when they submit a transaction. The reason for the lower throughput for the blockchain is believed to be the process that has to be carried out to add blocks to the blockchain. The reason the tangle has the lowest throughput is believed to be the long process that goes into deciding how to add new transactions to the directed acyclic graph (DAG).

If users wish to prioritize latency, the best option is the blockchain. The low latency for the blockchain is believed arise from the fact that nodes can accept the longest proof-of-work chain as the latest state of the system. The higher latency for the block-lattice is believed to arise from the fact that both a receive transaction and a send transaction have to be broadcast for a transaction to be completed. The reason the tangle has the highest latency is also believed to be the long process that goes into deciding how to add new transactions to the DAG.

Users may choose to use the block-lattice in a setting which requires high throughput such as a machine-to-machine micropayment system in the Internet-of-Things industry. Users may choose to use the blockchain in a setting which requires low latency such as real-life payments systems, where users don't want to wait a long time for a transaction to be finalized and become irreversible.

B. Importance of Results of Qualitative Analysis

The qualitative analysis shows that the tangle and the block-lattice guarantee the five fundamental properties defined by Xu et al. for blockchain-based systems [3]. This analysis shows that there is skepticism regarding the claim of equal rights for the participants in a blockchain system, which stems from the fact that a blockchain network relies on miners to aggregate valid transactions into blocks and append them to the blockchain. The tangle and the block-lattice guarantee equal rights for participants by making the users who submit a transaction participate in the validation of previous transactions, in the case of the tangle, or validation of their own transaction, in the case of the block-lattice.

Users may choose to use the tangle or the block-lattice in a permission-less public setting where all users should be able to join the network and submit and validate transactions. Users

may choose to use the blockchain in a permissioned setting where one or more authorities act as a gate for participation and not all users expect to have the same permissions as others. Blockchains may be more suitable for regulated industries such as banks for example, where permissions may include permission to join the network, permission to initiate transactions, and permissions to validate transactions.

C. Threats to Validity

The implementations of the data structures developed for this thesis do not precisely reflect all of the intricacies of the actual implementations of the Bitcoin blockchain, the IOTA tangle, or the Nano block-lattice. However, this thesis does represent a step forward in the study and comparison of different alternatives for underlying data structures of distributed ledgers because it captures the essence of three data structures by including by analyzing their fundamental functioning.

As mentioned before, in order to enable a comparison, the data structures implemented in this thesis are comprised of items which record one transaction each. This marks a clear difference from the way in which the Bitcoin blockchain is implemented since the blocks in that system hold the data for various transactions. For this reason, the throughput and latency data reported for the blockchain implemented for this thesis do not reflect the values for these quality attributes for the Bitcoin Blockchain. The Bitcoin blockchain has a theoretical maximum of 7 tps and a latency limited by the block generation time of 10 minutes. In a comparison with the data structures that were implemented, these would be the worst transaction processing metrics.

IX. LESSONS LEARNED

As explained before, to measure throughput and latency one node submits transactions to the network. We also conducted experiments in which multiple nodes submit transactions to the network. This experiment failed for the blockchain and the tangle because nodes only rewrite these data structures if the incoming data structure is larger than the one in memory, and this isn't guaranteed in a scenario in which multiple nodes are submitting transactions at the same time. This experiment succeeded for the block-lattice because nodes can update their data structure by appending incoming transactions to the relevant account-chains regardless of the size of the account-chains. This is an advantage of the block-lattice because, in its case, the size of the data structure is not a blocking factor when processing transactions.

X. FUTURE WORK

The implementations of the data structures developed for this thesis can be refined to better reflect the complexity of the Bitcoin blockchain, the IOTA tangle, and the Nano block-lattice. This may allow a more robust comparison of these data structures to be conducted.

Another analysis that could be conducted is testing whether a relaxation of some of the five fundamental properties could lead to an improvement in the technical challenges identified for the blockchain. As an example, Xu et al. discuss transparency in their analysis of different blockchain configurations, "using a public blockchain results in better

information transparency and auditability, but sacrifices performance” [3, p. 248]. Another fundamental property that could be adjusted to improve some of the quality attributes of distributed ledgers is equal rights. This fundamental property could be adjusted by setting up permissioned ledgers where one or more authorities act as a gate for participation. These authorities can decide who has permission to join the network, who has permission to initiate transactions, and who has permission to mine. Xu et al. explain that “there are often trade-offs between permissioned and permission-less blockchains including transaction processing rate, cost, censorship-resistance, reversibility, finality, and the flexibility in changing and optimizing network rules” [3, p. 245].

There are other options of underlying data structures that are being proposed and it would be interesting to compare them as well. Among these are the Hedera Hashgraph [24] proposed by professor Leemon Baird and the Avalanche Consensus Protocol [25] promoted by professor Emin Gün Sirer of Cornell University. These data structures are not included in this evaluation because they have a higher complexity and their implementation requires more time/resources than the ones available to complete this thesis.

XI. CONCLUSIONS

Research on distributed ledger technology has risen in the recent past and various alternatives of underlying structures for this technology have been proposed. The aim of this thesis is to address the optimality of different alternatives, by determining if and how they guarantee the fundamental properties of blockchains and if they can respond to some of the technical challenges identified for blockchains.

Quantitative and qualitative analyses allowed us to determine which data structures would be more suited for use in different real-life scenarios. It will be some time before distributed ledger technology is adopted for wide spread use, but this research represents a step in the right direction when it comes to paving the way to making this a reality.

REFERENCES

- [1] UK Government - Office for Science, “Distributed Ledger Technology: beyond block chain,” London, 2016.
- [2] S. Nakamoto, “Bitcoin: A peer-to-peer electronic cash system.” 2009.
- [3] X. Xu *et al.*, “A Taxonomy of Blockchain-Based Systems for Architecture Design,” in *2017 IEEE International Conference on Software Architecture (ICSA)*, 2017, pp. 243–252.
- [4] S. Popov, “The Tangle,” 2018.
- [5] C. Lemahieu, “Nano: A Feeless Distributed Cryptocurrency Network,” *White Pap.*, 2018.
- [6] M. Swan, *Blockchain: blueprint for a new economy*. Sebastopol, Calif.: O’Reilly Media, 2015.
- [7] K. Petersen, R. Feldt, S. Mujtaba, and M. Mattsson, “Systematic Mapping Studies in Software Engineering,” in *Proceedings of the 12th International Conference on Evaluation and Assessment in Software Engineering*, 2008, pp. 68–77.
- [8] B. Kitchenham, “Procedures for Performing Systematic Reviews,” Department of Computer Science, Keele University, UK, 2004.
- [9] J. Yli-Huomo, D. Ko, S. Choi, S. Park, and K. Smolander, “Where Is Current Research on Blockchain Technology?—A Systematic Review,” *PLoS One*, vol. 11, no. 10, pp. 1–27, 2016.
- [10] M. Dabbagh, M. Sookhak, and N. Safa, “The Evolution of Blockchain: A Bibliometric Study,” *IEEE Access*, vol. PP, p. 1, 2019.
- [11] P. Rimba, A. B. Tran, I. Weber, M. Staples, A. Ponomarev, and X. Xu, “Quantifying the Cost of Distrust: Comparing Blockchain and Cloud Services for Business Process Execution,” *Inf. Syst. Front.*, Aug. 2018.
- [12] I. Singh and S.-W. Lee, “Comparative Requirements Analysis for the Feasibility of Blockchain for Secure Cloud,” in *Requirements Engineering for Internet of Things*, 2018, pp. 57–72.
- [13] T. T. A. Dinh, R. Liu, M. Zhang, G. Chen, B. C. Ooi, and J. Wang, “Untangling Blockchain: A Data Processing View of Blockchain Systems,” *IEEE Trans. Knowl. Data Eng.*, vol. 30, no. 7, pp. 1366–1385, Jul. 2018.
- [14] Stack Overflow, “Stack Overflow’s Annual Developer Survey for 2019,” 2019.
- [15] D. Mingxiao, M. Xiaofeng, Z. Zhe, W. Xiangwei, and C. Qijun, “A review on consensus algorithm of blockchain,” in *2017 IEEE International Conference on Systems, Man, and Cybernetics (SMC)*, 2017, pp. 2567–2572.
- [16] The Iota Team, “Tangle visualization.” [Online]. Available: <https://github.com/iotaledger/iotavisualization>.
- [17] “The Go Programming Language - Documentation.” [Online]. Available: <https://golang.org/doc/>.
- [18] “A Tour of Go - Goroutines.” [Online]. Available: <https://tour.golang.org/concurrency/1>.
- [19] “A Tour of Go - Channels.” [Online]. Available: <https://tour.golang.org/concurrency/2>.
- [20] The Perlin Team, “The Noise Book.” [Online]. Available: <https://perlin-network.github.io/noise/noise.html>.
- [21] The Perlin Team, “Perlin.” [Online]. Available: <https://www.perlin.net/>.
- [22] Coral Health, “Code a simple P2P blockchain in Go!,” *Medium*, 2018. [Online]. Available: <https://medium.com/@mycoralhealth/code-a-simple-p2p-blockchain-in-go-46662601f417>.
- [23] The Nano Team, “About Nano.” [Online]. Available: <https://nano.org/en/about/>.
- [24] L. Baird, “The SWIRLDS Hashgraph Consensus Algorithm: Fair, Fast, Byzantine Fault Tolerance,” 2016.
- [25] Team Rocket, “Snowflake to Avalanche: A Novel Metastable Consensus Protocol Family for Cryptocurrencies,” 2018.

XII. ANNEXES

A. *Throughput Data*

TABLE II. THROUGHPUT DATA FOR THE BLOCKCHAIN. IN TRANSACTIONS/SECOND.

Data Structure	Run #	Number of Nodes		
		2 Nodes	4 Nodes	8 Nodes
Blockchain	Run 1	456	418	405
	Run 2	455	440	396
	Run 3	453	440	403
	Run 4	450	435	394
	Run 5	457	437	394
	Average	454	434	398
	Median	455	437	396
	StdDev	3	9	5

TABLE III. THROUGHPUT DATA FOR THE TANGLE. IN TRANSACTIONS/SECOND.

Data Structure	Run #	Number of Nodes		
		2 Nodes	4 Nodes	8 Nodes
Tangle	Run 1	72	78	74
	Run 2	65	80	72
	Run 3	80	73	76
	Run 4	73	68	60
	Run 5	82	73	82
	Average	74	74	73
	Median	73	73	74
	StdDev	7	5	8

TABLE IV. THROUGHPUT DATA FOR THE BLOCK-LATTICE. IN TRANSACTIONS/SECOND.

Data Structure	Run #	Number of Nodes		
		2 Nodes	4 Nodes	8 Nodes
Block-lattice	Run 1	6,197	4,216	3,933
	Run 2	6,061	4,109	3,920
	Run 3	6,208	3,603	3,788
	Run 4	6,071	4,020	3,810
	Run 5	6,005	4,470	3,892
	Average	6,108	4,084	3,869
	Median	6,071	4,109	3,892
	StdDev	90	317	66

B. Latency Data

1) 8 Nodes

TABLE V. LATENCY DATA FOR THE BLOCKCHAIN. 8 NODES. IN MILLISECONDS.

Data Structure	Node #	Transaction #				
		#50	#100	#200	#400	#800
Blockchain	Node 2	6	4	4	9	14
	Node 3	6	5	7	8	22
	Node 4	5	4	6	8	21
	Node 5	4	5	7	10	23
	Node 6	8	5	6	10	24
	Node 7	3	5	6	12	24
	Node 8	6	6	7	6	31
	Average	6	5	6	9	23
	Median	6	5	6	9	23
	StdDev	1	1	1	2	5

TABLE VI. LATENCY DATA FOR THE TANGLE. 8 NODES. IN MILLISECONDS.

Data Structure	Node #	Transaction #
		#50
Tangle	Node 2	5
	Node 3	5
	Node 4	5
	Node 5	5
	Node 6	4
	Node 7	4
	Node 8	5
	Average	5
	Median	5
	StdDev	0

TABLE VII. LATENCY DATA FOR THE BLOCK-LATTICE. 8 NODES. IN MILLISECONDS.

Data Structure	Node #	Transaction #				
		#50	#100	#200	#400	#800
Block-lattice	Node 1	57	42	69	107	150
	Node 3	55	41	70	128	150
	Node 4	56	41	68	155	265
	Node 5	59	45	71	133	193
	Node 6	59	45	61	120	280
	Node 7	58	44	65	152	150
	Node 8	56	42	66	147	130
	Average	57	43	67	135	188
	Median	57	42	68	133	150
	StdDev	2	2	4	18	61

2) 4 Nodes

TABLE VIII. LATENCY DATA FOR THE BLOCKCHAIN. 4 NODES. IN MILLISECONDS.

Data Structure	Node #	Transaction #				
		#50	#100	#200	#400	#800
Blockchain	Node 2	14	7	8	8	9
	Node 3	16	5	10	9	16
	Node 4	14	4	6	9	19
	Average	15	5	8	9	15
	Median	14	5	8	9	16
	StdDev	1	1	2	0	5

TABLE IX. LATENCY DATA FOR THE TANGLE. 4 NODES. IN MILLISECONDS.

Data Structure	Node #	Transaction #
		#50
Tangle	Node 2	3
	Node 3	4
	Node 4	4
	Average	4
	Median	4
	StdDev	0

TABLE X. LATENCY DATA FOR THE BLOCK-LATTICE. 4 NODES. IN MILLISECONDS.

Data Structure	Node #	Transaction #				
		#50	#100	#200	#400	#800
Block-lattice	Node 1	65	46	67	124	263
	Node 3	63	44	65	148	279
	Node 4	64	44	65	182	353
	Average	64	45	66	151	299
	Median	64	44	65	148	279
	StdDev	1	1	1	29	48

3) 2 Nodes

TABLE XI. LATENCY DATA FOR THE DATA STRUCTURES. 2 NODES. IN MILLISECOND

Data Structure	Node #	Transaction #				
		#50	#100	#200	#400	#800
Blockchain	Node 2	6	5	8	13	12
Tangle	Node 2	8				
Block-lattice	Node 1	64	36	67	110	243